MASSACHVSETTS INSTITVTE OF TECHNOLOGY
Department of Electrical Engineering and Computer Science

6.945/6.S081 Spring 2014
Problem Set 1

Issued: Wed. 12 Feb. 2014   Due: Wed. 19 Feb. 2014

Readings:
    Review SICP chapter 1 (especially section 1.3)
    MIT/GNU Scheme documentation, sections 5 and 6
      (characters and strings)
    Debian GNU/Linux info on regular expressions
      from the grep man page (attached). This is sane.

Code:
    regexp.scm, tests.txt (both attached)
    Windows grep: http://gnuwin32.sourceforge.net/packages/grep.htm

Documentation:
  The MIT/GNU Scheme installation and documentation can
    be found online at http://www.gnu.org/software/mit-scheme/
    The reference manual is in:
  http://www.gnu.org/software/mit-scheme/documentation/mit-scheme-ref/

    The (insane) POSIX manual page for regular expressions:
  http://www.opengroup.org/onlinepubs/009695399/basedefs/xbd_chap09.html

Regular Expressions

Regular expressions are ubiquitous. On the surface, regular
expressions look like a combinator language, because expression
fragments can be combined to make more complex expressions. But the
meaning of a fragment is highly dependent on the expression it is
embedded in. For example, to include a caret character in a bracket
expression, [...], it must not be in the first character position.
If the caret appears after the first character it is just an ordinary
character, but if it appears as the first character it negates the
meaning of the bracket expression. For example, a bracket expression
may not contain just a caret.

So the syntax of the regular-expression language is awful. There are
various incompatable forms of the language and the quotation
conventions are baroquen [sic]. Nevertheless, there is a great deal
of useful software, for example grep, that uses regular expressions to
specify the desired behavior.

Although regular-expression systems are derived from a perfectly good
mathematical formalism, the particular choices made by implementers to
expand the formalism into useful software systems are often
disastrous: the quotation conventions adopted are highly irregular;
the egregious misuse of parentheses, both for grouping and for
backward reference, is a miracle to behold. In addition, attempts to
increase the expressive power and address shortcomings of earlier
designs have led to a proliferation of incompatible derivative
languages.

Part of the value of this problem set is to experience how bad things
can be. Here we will invent both a combinator language for specifying
regular expressions and a means of translating this language to
conventional regular-expression syntax. The application is to be able
to use the capabilities of systems like grep from inside the Scheme
environment. This will give us all the advantages of a combinator
language. It will have a clean, modular description while retaining
the ability to use existing tools. Users of this language will have
nothing to grep about, unless they value concise expression over
readability.

As with any language there are primitives, means of combination, and
means of abstraction. Our language allows the construction of
patterns that utilities like grep can match against character-string
data. Because this language is embedded in Scheme we inherit all of
the power of Scheme: we can use Scheme constructs to combine patterns
and Scheme procedures to abstract them.

                    A regular expression combinator language

Patterns are built out of primitive patterns. The primitive pattern
elements are:

  (r:dot) matches any character except newline

  (r:bol) matches only the beginning of a line

  (r:eol) matches only the end of a line

  (r:quote <string>) matches the given string

  (r:char-from <char-set>)
      matches one character that is in the given string

  (r:char-not-from <char-set>)
      matches one character that is not in the given string

Patterns can be combined to make compound patterns:

```
(r:seq <pattern> ...)
     This pattern matches each of the argument patterns in sequence,
     from left to right

(r:alt <pattern> ...)
     This pattern tries each argument pattern from left to right,
     until one of these alternatives matches.  If none matches then
     this pattern does not match.

(r:repeat <min> <max> <pattern>)
     This pattern tries to match the given argument pattern a minimum
     of min times but no more than a maximum of max times.  If max is
     given as #f then there is no maximum specified.  Note that if
     max=min the given pattern must be matched exactly that many
     times.
```

Because these are all Scheme procedures (in the file regexp.scm) you
can freely mix these with any Scheme code.

Here are some examples:

```
Pattern:  (r:seq (r:quote "a") (r:dot) (r:quote "c"))
Matches:  "abc" and "aac" and "acc"

Pattern:  (r:alt (r:quote "foo") (r:quote "bar") (r:quote "baz"))
Matches:  "foo" and "bar" and "baz"

Pattern:  (r:repeat 3 5 (r:alt (r:quote "cat") (r:quote "dog")))
Matches:  "catdogcat" and "catcatdogdog" and "dogdogcatdogdog"
          but not "catcatcatdogdogdog"

Pattern:  (let ((digit
                 (r:char-from "0123456789")))
            (r:seq (r:bol)
                   (r:quote "[")
                   digit
                   digit
                   (r:quote "]")
                   (r:quote ".")
                   (r:quote " ")
                   (r:char-from "ab")
                   (r:repeat 3 5 (r:alt (r:quote "cat") (r:quote "dog")))
                   (r:char-not-from "def")
                   (r:eol)))
Matches:  "[09]. acatdogdogcats" but not
          "[10]. ifacatdogdogs" nor
          "[11]. acatdogdogsme"
```

In the file regexp.scm we define an interface to the grep utility,
which allows a Scheme program to call grep with a regular expression
(constructed from the given pattern) on a given file name. The grep
utility extracts lines that contain a substring that can match the
given pattern.

So, for example: given the test file test.txt (supplied) we can find
the lines in the file that contain a match to the given pattern.

```
(pp
 (r:grep (r:seq " "
                (r:repeat 3 5 (r:alt (r:quote "cat") (r:quote "dog")))
                (r:eol))
          "tests.txt"))
("[09]. catdogcat" "[10]. catcatdogdog" "[11]. dogdogcatdogdog")
;Unspecified return value
```

Note that the pretty-printer (pp) returns an unspecified value, after
printing the list of lines that grep found matching the pattern
specified.

                        Implementation

Let's look at how this language is implemented. Regular expressions
will be represented as strings using the "basic" regular expression
syntax of the Unix grep program.

```
(define (r:dot) ".")
(define (r:bol) "^")
(define (r:eol) "$")
```

These directly correspond to regular-expression syntax. Next, r:seq
implements a way to treat a given set of regular-expression fragments
as a self-contained element:

```
(define (r:seq . exprs)
  (string-append "\\(" (apply string-append exprs) "\\)"))
```

The use of parentheses in the result isolates the content of the given
expression fragments from the surrounding context. The implementation
of r:quote is a bit harder:

```
(define (r:quote string)
  (r:seq
   (list->string
    (append-map (lambda (char)
                  (if (memv char chars-needing-quoting)
                      (list #\\ char)
                      (list char)))
                (string->list string)))))
```

```
(define chars-needing-quoting
  '(#\. #\[ #\\ #\^ #\$ #\*))
```

In a regular expression, most characters are self-quoting. However,
some characters are regular-expression operators and must be

explicitly quoted. We wrap the result using r:seq to guarantee that
the quoted string is self contained.

To implement alternative subexpressions, we interpolate a vertical bar
between subexpressions and wrap the result using r:seq:

```
(define (r:alt . exprs)
  (if (pair? exprs)
      (apply r:seq
             (cons (car exprs)
                   (append-map (lambda (expr)
                                 (list "\\|" expr))
                               (cdr exprs))))
      (r:seq)))
```

Note that alternative expressions, unlike the rest of the regular
expressions supported here, are not a part of POSIX Basic Regular
Expression (BRE) syntax. They are an extension defined by GNU grep,
which is supported by many implementations.

The implementation of repetition is straightforward by using copies of
the given regular expression:

```
(define (r:repeat min max expr)
  (apply r:seq
         (append (make-list min expr)
                 (if (eqv? max min)
                     '()
                     (if max
                         (make-list (- max min)
                                    (r:alt expr ""))
                         (list expr "*"))))))
```

This makes min copies, followed by (- max min) optional copies. (Each
optional copy is an alternative of the expression and an empty
expression.) If there's no maximum, then the expression followed by
an asterisk matches any number of times.

The implementation of r:char-from and r:char-not-from is complicated
by the need for baroque quotation. This is best organized in two
parts, the first to handle the differences between them, and the
second for the common quotation:

```
(define (r:char-from string)
  (case (string-length string)
    ((0) (r:seq))
    ((1) (r:quote string))
    (else
     (bracket string
              (lambda (members)
                (if (lset= eqv? '(#\- #\^) members)
                    '(#\- #\^)
                    (quote-bracketed-contents members)))))))

(define (r:char-not-from string)
  (bracket string
           (lambda (members)
             (cons #\^ (quote-bracketed-contents members)))))
```

```
(define (bracket string procedure)
  (list->string
   (append '(#\[)
           (procedure (string->list string))
           '(#\]))))
```

The special cases for r:char-from handle empty and singleton
sets of characters specially, which simplifies the general case.
There is also a particular special case for a set containing only
caret and hyphen. But r:char-not-from has no such
restrictions. The general case handles the three characters that have
special meaning inside a bracket by placing them in positions where
they are not operators. (We told you this was ugly!)

```
(define (quote-bracketed-contents members)
  (let ((optional
         (lambda (char)
           (if (memv char members) (list char) '()))))
    (append (optional #\])
            (remove (lambda (c)
                      (memv c chars-needing-quoting-in-brackets))
                    members)
            (optional #\^)
            (optional #\-))))
```

```
(define chars-needing-quoting-in-brackets
  '(#\] #\^ #\-))
```

In order to test this code, we can print the corresponding grep
command and use cut and paste to run it in a shell:

```
(define (write-bourne-shell-grep-command expr filename)
  (display (bourne-shell-grep-command-string expr filename)))
```

```
(define (bourne-shell-grep-command-string expr filename)
  (string-append "grep -e "
                 (bourne-shell-quote-string expr)
                 " "
                 filename))
```

Because shells have their own quoting issues, we need to not only
quote the regular expression, but also choose which shell to use,
since different shells use different quoting conventions. The Bourne
shell is ubiquitous, and has a relatively simple quoting convention:

```
(define (bourne-shell-quote-string string)
  (list->string
   (append (list #\')
           (append-map (lambda (char)
                         (if (char=? char #\')
                             (list #\' #\\ char #\')
                             (list char)))
                       (string->list string))
           (list #\'))))
```

This quoting convention uses single-quote characters surrounding a
string, which quotes anything in the string other than a single-quote,

which ends the quoted string. So, to quote a single-quote character,
we must end the string, quote the single quote explicitly using
backslash, and then start another quoted string. The shell interprets
this concatenation as a single token. (Are we having fun yet?)

We can directly use the system grep from MIT/GNU Scheme. (This code
is MIT/GNU Scheme specific.)

```scheme
(load-option 'synchronous-subprocess)

(define (r:grep expr filename)
  (let ((port (open-output-string)))
    (and (= (run-shell-command
              (bourne-shell-grep-command-string expr filename)
              'output port)
            0)
         (r:split-lines (get-output-string port)))))

(define (r:split-lines string)
  (reverse
   (let ((end (string-length string)))
     (let loop ((i 0) (lines '()))
       (if (< i end)
           (let ((j
                   (substring-find-next-char string i end #\newline)))
             (if j
                 (loop (+ j 1)
                       (cons (substring string i j) lines))
                 (cons (substring string i end) lines)))
           lines)))))
```

However, we can avoid most of this quotation hair by using the
run-synchronous-subprocess feature of MIT/GNU Scheme to call the grep
program without the intermediate complication of a shell. To do this
we can use:

```
(define (r:grep* expr filename)
  (let ((port (open-output-string)))
    (run-synchronous-subprocess "grep"
                                (list expr filename)
                                'output
                                port)
    (r:split-lines (get-output-string port))))
```

The moral of this story

Our translator is very complicated because most regular expressions
are not composable to make larger regular expressions unless extreme
measures are taken to isolate the parts. Our translator does this
work, but consequently the regular expressions that it generates have
much unnecessary boilerplate. Humans don't write regular expressions
this way because they use boilerplate only where necessary---but often
miss instances where it is necessary, causing hard-to-find bugs.

The moral of this story is that regular expressions are a beautiful
example of how not to build a system. Using composable parts and
combinators to make new parts by combining others leads to simpler and
more robust implementations.

Problem 1.1: Warmup

In the traditional regular expression language the asterisk (*)
operator following a subpattern means zero or more copies of the
subpattern and the plus-sign (+) operator following a subpattern means
one or more copies of the subpattern. Define Scheme procedures r:*
and r:+ to take a pattern and iterate it as necessary. This can be
done in terms of r:repeat.

Demonstrate your procedures on real data in complex patterns.

          A Subtle Bug, One Bad Joke, Two Tweaks, and a Revelation

Ben Bitdiddle has noticed a problem with our implementation of
(r:repeat <min> <max> <pattern>).

The use of (r:alt expr "") at the end of the r:repeat procedure is
a bit dodgy. This code fragment compiles to an Extended Regular
Expression (ERE) Alternation regular expression of the form (expr|).
(See 9.4.7 of the POSIX regular expression document.)

This relies on the fact that alternation with something and nothing is
the equivalent of saying "one or none". That is: (expr|) denotes one
or no instances of expr. Unfortunately, this depends on an
undocumented GNU extension to the formal POSIX standard for REs.

Specifically, section 9.4.3 states that a vertical line appearing
immediately before a close parenthesis (or immediately after an open
parenthesis) produces undefined behavior. In essence, an RE must not
be a null sequence.

GNU grep just happens to Do The Right Thing (tm) when presented with
(x|). Not all grep implementations are as tolerant.

Therefore, Ben asks his team of three code hackers (Louis, Alyssa and his niece Bonnie) to propose alternative workarounds. Ultimately, he proposes his own patch, which you will implement.

  * Louis Reasoner suggests that a simple, elegant fix would be to
    replace the code fragment (r:alt expr "") with a straight-
    forward call to (r:repeat 0 1 expr).

  * Alyssa P. Hacker proposes that an alternative fix would be to
    rewrite the else clause of r:repeat to compile (r:repeat 3 5 <x>)
    into the equivalent of (xxx|xxxx|xxxxx) instead of the naughty
    xxx(x|)(x|) non-POSIX-compliant undefined regular expression. She
    refers to section 9.4.7 of the POSIX regular expression document.

  * Bonnie Bitdiddle points to the question mark (?) operator in
    section 9.4.6.4 and proposes that a better fix would be to
    implement an r:? operator then replace (r:alt expr "") with
    (r:? expr).

  * Meanwhile, Ben looks closely at the RE spec and has a revelation.
    He proposes that r:repeat be re-implemented to emit Interval
    expressions. See section 9.3.6.5 of the POSIX documentation.
    Please try not to get sick.

Problem 1.2: The Proposals

Let's very briefly consider each proposal:

a. Everyone in the room immediately giggles at Louis's silly joke.
   What's so funny about it? That is, what's wrong with this idea?

   A one-sentence punchline will do.

b. What advantages does Bonnie's proposal have over Alyssa's
   in terms of both code and data?

   A brief, concise yet convincing few sentences suffice.

c. What advantage does Ben's proposal have over all the others?
   Specifically, ponder which section of the POSIX document he cites
   versus which sections the others cite, then take a quick peek at
   Problem 1.5 below and consider the implications. Also, consider
   the size of the output strings in this new code as well as the
   overall clarity of the code.

   Again, a brief sentence or two is sufficient.

d. Following Ben's proposal, re-implement r:repeat to emit Interval
   Expressions. Hint: Scheme's number->string procedure should be
   handy. Caveat: Beware the backslashes.

   Show the output it generates on a few well-chosen sample inputs.
   Demonstrate your procedure on real data in complex patterns.

Too Much Nesting

Our program produces excessively nested regular expressions: it makes
groups even when they are not necessary. For example, the following
simple pattern leads to an overly complex regular expression:

```
(display (r:seq (r:quote "a") (r:dot) (r:quote "c")))
\(\(a\).\(c\)\)
```

Another problem is that BREs may involve back-references. (See
9.3.6.3 of the POSIX regular expression documentation.) A
back-reference refers to a previously parenthesized subexpression. So
it is important that the parenthesized subexpressions be ones
explicitly placed by the author of the pattern. (Aargh! This is one
of the worst ideas I (GJS) have ever heard of -- grouping, which is
necessary for iteration, was confused with naming for later reference.
What a crock!)

Problem 1.3: Optimization

Edit our program to eliminate as much of the unnecessary nesting as
you can. Caution: there are subtle cases here that you have to watch
out for. What is such a case? Demonstrate your better version of our
program and show how it handles the subtleties.

Hint: Our program uses strings as its intermediate representation as
well as its result. You might consider using a different intermediate
representation.

Problem 1.4: Back-references

Add in a procedure for constructing back-references.
Have fun getting confused about BREs.

Standards?

> The best thing about standards is that
> there are so many to choose from.

There are also Extended Regular Expressions (EREs) defined in the
POSIX regular expression documentation. Some software, such as egrep,
uses this version of regular expressions. Unfortunately EREs are not
a conservative extension of BREs: ERE syntax is actually inconsistent
with BRE syntax! It is an interesting project to extend our Scheme
pattern language so that the target can be either BREs or EREs.

Problem 1.5: Ugh!

  a. What are the significant differences between BREs and EREs that
  make this a pain? List the differences that must be addressed.

  b. How can the back end be factored so that our language can compile
  into either kind of regular expression, depending on what is needed?
  How can we maintain the abstract layer that is independent of the
  target regular expression language? Explain your strategy.

  c. Extend our implementation to have both back ends.

Demonstrate your work by making sure that you can run egrep as well as
grep, with equivalent results in cases that test the differences you
found in part a.

       End of Problem Set. Reference Material Follows.

The following is an excerpt from the Debian GNU/Linux man page on grep.

REGULAR EXPRESSIONS
     A regular expression is a pattern that describes a set of strings.
     Regular expressions are constructed analogously to arithmetic expres-
     sions, by using various operators to combine smaller expressions.

     Grep understands three different versions of regular expression syn-
     tax: "basic," "extended," and "perl." In GNU grep, there is no dif-
     ference in available functionality using either of the first two syn-
     taxes. In other implementations, basic regular expressions are less
     powerful.   The following description applies to extended regular
     expressions; differences for basic regular expressions are summarized
     afterwards.   Perl regular expressions add additional functionality,
     but the implementation used here is undocumented and is not compati-
     ble with other grep implementations.

     The fundamental building blocks are the regular expressions that
     match a single character. Most characters, including all letters and
     digits, are regular expressions that match themselves. Any metachar-
     acter with special meaning may be quoted by preceding it with a back-
     slash.

     A bracket expression is a list of characters enclosed by [and]. It
     matches any single character in that list; if the first character of
     the list is the caret ^ then it matches any character not in the
     list. For example, the regular expression [0123456789] matches any
     single digit.

     Within a bracket expression, a range expression consists of two char-
     acters separated by a hyphen. It matches any single character that
     sorts between the two characters, inclusive, using the locale's col-
     lating sequence and character set. For example, in the default C
     locale, [a-d] is equivalent to [abcd]. Many locales sort characters
     in dictionary order, and in these locales [a-d] is typically not
     equivalent to [abcd]; it might be equivalent to [aBbCcDd], for exam-
     ple. To obtain the traditional interpretation of bracket expres-
     sions, you can use the C locale by setting the LC_ALL environment
     variable to the value C.

     Finally, certain named classes of characters are predefined within
     bracket expressions, as follows. Their names are self explanatory,
     and they are [:alnum:], [:alpha:], [:cntrl:], [:digit:], [:graph:],
     [:lower:],   [:print:],   [:punct:],   [:space:],   [:upper:],   and
     [:xdigit:]. For example, [[:alnum:]] means [0-9A-Za-z], except the
     latter form depends upon the C locale and the ASCII character encod-
     ing, whereas the former is independent of locale and character set.
     (Note that the brackets in these class names are part of the symbolic
     names, and must be included in addition to the brackets delimiting
     the bracket list.)   Most metacharacters lose their special meaning
     inside lists. To include a literal] place it first in the list.
     Similarly, to include a literal ^ place it anywhere but first.
     Finally, to include a literal - place it last.

The period . matches any single character. The symbol \w is a synonym for [[:alnum:]] and \W is a synonym for [^[:alnum]].

The caret ^ and the dollar sign $ are metacharacters that respectively match the empty string at the beginning and end of a line. The symbols \< and \> respectively match the empty string at the beginning and end of a word. The symbol \b matches the empty string at the edge of a word, and \B matches the empty string provided it's not at the edge of a word.

A regular expression may be followed by one of several repetition operators:
```
?       The preceding item is optional and matched at most once.
*       The preceding item will be matched zero or more times.
+       The preceding item will be matched one or more times.
{n}     The preceding item is matched exactly n times.
{n,}    The preceding item is matched n or more times.
{n,m}   The preceding item is matched at least n times, but not more
        than m times.
```

Two regular expressions may be concatenated; the resulting regular expression matches any string formed by concatenating two substrings that respectively match the concatenated subexpressions.

Two regular expressions may be joined by the infix operator |; the resulting regular expression matches any string matching either subexpression.

Repetition takes precedence over concatenation, which in turn takes precedence over alternation. A whole subexpression may be enclosed in parentheses to override these precedence rules.

The backreference \n, where n is a single digit, matches the substring previously matched by the nth parenthesized subexpression of the regular expression.

In basic regular expressions the metacharacters ?, +, {, |, (, and) lose their special meaning; instead use the backslashed versions \?, \+, \{, \|, \(, and \).

Traditional egrep did not support the { metacharacter, and some egrep implementations support \{ instead, so portable scripts should avoid { in egrep patterns and should use [{] to match a literal {.

GNU egrep attempts to support traditional usage by assuming that { is not special if it would be the start of an invalid interval specification. For example, the shell command egrep '{1' searches for the two-character string {1 instead of reporting a syntax error in the regular expression. POSIX.2 allows this behavior as an extension, but portable scripts should avoid it.

```scheme
;;;; Scheme Regular Expression Language Implementation -- regexp.scm

(define (r:dot) ".")
(define (r:bol) "^")
(define (r:eol) "$")

(define (r:quote string)
  (r:seq
   (list->string
    (append-map (lambda (char)
                  (if (memv char chars-needing-quoting)
                      (list #\\ char)
                      (list char)))
                (string->list string)))))

(define chars-needing-quoting
  '(#\. #\[ #\\ #\^ #\$ #\*))

(define (r:char-from string)
  (case (string-length string)
    ((0) (r:seq))
    ((1) (r:quote string))
    (else
     (bracket string
              (lambda (members)
                (if (lset= eqv? '(#\- #\^) members)
                    '(#\- #\^)
                    (quote-bracketed-contents members)))))))

(define (r:char-not-from string)
  (bracket string
           (lambda (members)
             (cons #\^ (quote-bracketed-contents members)))))

(define (bracket string procedure)
  (list->string
   (append '(#\[)
           (procedure (string->list string))
           '(#\]))))

(define (quote-bracketed-contents members)
  (let ((optional
         (lambda (char) (if (memv char members) (list char) '()))))
    (append (optional #\])
            (remove (lambda (c)
                      (memv c chars-needing-quoting-in-brackets))
                    members)
            (optional #\^)
            (optional #\-))))

(define chars-needing-quoting-in-brackets
  '(#\] #\^ #\-))
```

```scheme
;;; Means of combination for patterns

(define (r:seq . exprs)
  (string-append "\\(" (apply string-append exprs) "\\)"))

;;; An extension to POSIX basic regular expressions.
;;; Supported by GNU grep and possibly others.
(define (r:alt . exprs)
  (if (pair? exprs)
      (apply r:seq
             (cons (car exprs)
                   (append-map (lambda (expr)
                                 (list "\\|" expr))
                               (cdr exprs))))
      (r:seq)))


(define (r:repeat min max expr)
  (apply r:seq
         (append (make-list min expr)
                 (if (eqv? max min)
                     '()
                     (if max
                         (make-list (- max min)
                                    (r:alt expr ""))
                         (list expr "*"))))))
```

```
;;; Using system's grep.
(define (write-bourne-shell-grep-command expr filename)
  (display (bourne-shell-grep-command-string expr filename)))

(define (bourne-shell-grep-command-string expr filename)
  (string-append "grep -e "
                 (bourne-shell-quote-string expr)
                 " "
                 filename))

;;; Works for any string without newlines.
(define (bourne-shell-quote-string string)
  (list->string
   (append (list #\')
           (append-map (lambda (char)
                         (if (char=? char #\')
                             (list #\' #\\ char #\')
                             (list char)))
                       (string->list string))
           (list #\'))))


;;; This is MIT/Scheme specific and compatible with grep for the
;;; purposes of this code.

(load-option 'synchronous-subprocess)

(define (r:grep expr filename)
  (let ((port (open-output-string)))
    (and (= (run-shell-command
             (bourne-shell-grep-command-string expr filename)
             'output port)
            0)
         (r:split-lines (get-output-string port)))))

(define (r:split-lines string)
  (reverse
   (let ((end (string-length string)))
     (let loop ((i 0) (lines '()))
       (if (< i end)
           (let ((j
                  (substring-find-next-char string i end #\newline)))
             (if j
                 (loop (+ j 1)
                       (cons (substring string i j) lines))
                 (cons (substring string i end) lines)))
           lines)))))
```

```
#|
;;; For example...

(pp (r:grep (r:seq (r:quote "a") (r:dot) (r:quote "c")) "tests.txt"))
("[00]. abc"
 "[01]. aac"
 "[02]. acc"
 "[03]. zzzaxcqqq"
 "[10]. catcatdogdog"
 "[12]. catcatcatdogdogdog")
;Unspecified return value

;;; And...

(pp (r:grep (r:alt (r:quote "foo") (r:quote "bar") (r:quote "baz"))
            "tests.txt"))
("[05]. foo" "[06]. bar" "[07]. foo bar baz quux")
;Unspecified return value


(pp (r:grep (r:repeat 3 5 (r:alt (r:quote "cat") (r:quote "dog")))
            "tests.txt"))
("[09]. catdogcat"
 "[10]. catcatdogdog"
 "[11]. dogdogcatdogdog"
 "[12]. catcatcatdogdogdog"
 "[13]. acatdogdogcats"
 "[14]. ifacatdogdogs"
 "[15]. acatdogdogsme")
;Unspecified return value

(pp
 (r:grep (r:seq " "
                (r:repeat 3 5 (r:alt (r:quote "cat") (r:quote "dog")))
                (r:eol))
         "tests.txt"))
("[09]. catdogcat" "[10]. catcatdogdog" "[11]. dogdogcatdogdog")
;Unspecified return value
```

```
(pp
 (r:grep
  (let ((digit
          (r:char-from "0123456789")))
    (r:seq (r:bol)
           (r:quote "[")
           digit
           digit
           (r:quote "]")
           (r:quote ".")
           (r:quote " ")
           (r:char-from "ab")
           (r:repeat 3 5 (r:alt "cat" "dog"))
           (r:char-not-from "def")
           (r:eol)))
   "tests.txt"))
("[13]. acatdogdogcats")
;Unspecified return value
|#
```

;;; This is the file tests.txt

[00]. abc
[01]. aac
[02]. acc
[03]. zzzaxcqqq
[04]. abdabec

[05]. foo
[06]. bar
[07]. foo bar baz quux
[08]. anything containing them

[09]. catdogcat
[10]. catcatdogdog
[11]. dogdogcatdogdog
[12]. catcatcatdogdogdog

[13]. acatdogdogcats
[14]. ifacatdogdogs
[15]. acatdogdogsme

;;; This is the file tests.txt