MASSACHVSETTS INSTITVTE OF TECHNOLOGY
Department of Electrical Engineering and Computer Science

6.945 Spring 2014
Problem Set 9

Issued: Wed. 16 April 2014   Due: Wed. 30 April 2014

Note: This is the last problem set for this term. You now have a term
      project to work on. Also, this problem set is optional for
      students enrolled in 6.S081.

Readings:

  Online MIT/GNU Scheme Documentation,
      Section 2.3: Dynamic Binding - fluid-let
      Section 12.4: Continuations   - call-with-current-continuation

  Here is a nice paper about continuations and threads:
      http://repository.readscheme.org/ftp/papers/sw2003/Threads.pdf

  In fact, there is an entire bibliography of stuff about this on:
      http://library.readscheme.org/page6.html

Code: load.scm ghelper.scm utils.scm time-share.scm schedule.scm
      syntax.scm rtdata.scm interp-actor.scm repl.scm

The Actor Model of Concurrency

We have examined traditional time-sharing, and the propagator system
we have played with is one way to think about multiple computing
agents acting in parallel: each propagator can be thought of as an
independent computing agent. But we didn't actually implement the
propagator system so that it could work in parallel. To make that
work we would have to arrange that cells could serialize update
requests so that they would not suffer from reader-writer problems.
Also, the propagator system, as implemented, manipulates worldviews by
global side-effects on the status of premises. In a concurrent system
it would be appropriate for the worldview to be represented separately
for each query process. Rebuilding the propagator system to work well
concurrently is an excellent project, but here we go down another path
for concurrency.

In 1973 Carl Hewitt came up with an idea for formulating concurrency,
based on message passing, called "actors." (Yes, that is the same
Carl Hewitt who had the idea of match combinators.) Very crudely,
Hewitt's idea is that when you call an actor (send it a message) the

call immediately returns to the caller, but the message is put on a
queue of work for the actor to do. Each actor processes the elements
of its input queue, perhaps changing its local state as it does so.
In Hewitt's actors the state update is accomplished by the actor
"replacing itself" with a new behavior. If the intention of the
caller was to receive a reply it would have to pass its return
continuation as an argument to the called actor as part of its calling
message.

In Hewitt's ideal system there were no computational objects that were
not actors, including integers. So to add "2" to "3" you send an
"add" message with a return address to "2". You then receive a
2-adder to which you send a "3" and a return address which will
receive "5". This made a very elegant object-oriented system, where
everything was implemented by message passing. See the article
http://en.wikipedia.org/wiki/Actor_model for more information.

Here I will be less orthodox and consider an actor-inspired system
that adds an actor-like procedure to an ordinary Scheme-like lambda-
calculus interpreter. This is more to my (GJS) liking, because it
does not require everything to be an actor -- I hate grand theories
and magic bullets that restrict a programmer to do things by someone's
grand theory. The languages Erlang and Scala both provide mechanisms
to support actor-like behaviors.

So, let's get into it. We provide an embedded actor interpreter
system -- you can load it with (load "load"). We initialize
this interpreter with (init) and if we get an error, we can reenter it
with (go). It provides the prompt "eval>".

This actor simulator is a work in progress. THERE ARE BUGS! However
there is much that we can learn from this system.

In this system we can run ordinary programs in the usual way:

```
    eval> (define fib
            (lambda (n)
              (if (< n 2)
                  n
                  (+ (fib (- n 1))
                     (fib (- n 2)))))))

    eval> (fib 10)
    55
```

But we can also define the Fibonacci computer in terms of actors,
using "alpha expressions" rather than "lambda expressions":

```
eval> (define fib
        (alpha (n c)
          (if (< n 2)
              (c n)
              (let ((x 'not-ready) (y 'not-ready))
                (define wait-for-xy
                  (alpha (k)
                         (if (boolean/or
                              (eq? x 'not-ready)
                              (eq? y 'not-ready))
                             (wait-for-xy k)
                             (k #t))))
                (fib (- n 1)
                     (lambda (v) (set! x v)))
                (fib (- n 2)
                     (lambda (v) (set! y v)))
                (wait-for-xy
                 (lambda (ignore) (c (+ x y)))))))))

eval> (fib 10 write-line)
55
done
```

Here fib is an actor that takes a number and a continuation to call
with the value. If the number is not less than 2 it sets up two state
variables, x and y, to capture the values of the recursive calls. It
also defines an actor procedure that busy-waits until x and y are
available. The continuations of the recursive calls are just
assignments to the capture variables. In cases where we don't want
to start a concurrent process, we use lambda rather than alpha.

                        Implementation

We will start with a simple interpreter to build an experimental actor
system. The changes to the interpreter system are quite minimal.
First, we add a special-form syntax for alpha expressions:

```
  (define (alpha? exp) (tagged-list? exp 'alpha))

  (define (alpha-parameters alpha-exp) (cadr alpha-exp))

  (define (alpha-body alpha-exp)
    (let ((full-body (cddr alpha-exp)))
      (sequence->begin full-body)))
```

We did not change or add to the syntax for definitions, so expressions
of the form

```
(define (foo x y) <body>)
```

define ordinary compound procedures with lambda. Perhaps we should
have a defactor?

Since eval is defined as generic we need to produce a dispatch for
alpha expressions that makes an actor procedure, analogous to a
lambda-specified procedure:

```
(defhandler eval
  (lambda (expression environment)
    (make-actor-procedure
     (alpha-parameters expression)
     (alpha-body expression)
     environment))
  alpha?)
```

Now, we need to specify the run-time data structures for such actor
procedures:

```
(define-record-type actor-procedure
    (make-actor-procedure% vars bproc env task-queue runnable)
    actor-procedure?
  (vars  actor-parameters)
  (bproc actor-body)
  (env   actor-environment)
  (task-queue get-actor-task-queue set-actor-task-queue!)
  (runnable actor-runnable? set-actor-runnable!))

(define (make-actor-procedure vars bproc env)
  (make-actor-procedure% vars bproc env (queue:make) #f))
```

Note that actor procedures differ from ordinary compound procedures
produced by lambda expressions, in that actor procedures have a task
queue and a bit that is true only if the actor is scheduled to run.

Next, we augment the apply generic operation to handle actor
procedures:

```
(defhandler apply
  (lambda (actor operands calling-environment)
    (if (not (= (length (actor-parameters actor))
                (length operands)))
        (error "Wrong number of operands supplied"))
    (let ((arguments
           (map (lambda (parameter operand)
                  (evaluate-procedure-operand
                   parameter operand calling-environment))
                (actor-parameters actor)
                operands)))
      (add-to-tasks! actor
                     (lambda ()
                       (eval (actor-body actor)
                             (extend-environment
                              (map procedure-parameter-name
                                   (actor-parameters actor))
                              arguments
                              (actor-environment actor)))))
      'actor-applied)
  actor-procedure?)
```

This is just like an ordinary compound procedure except that we defer
the actual work to be done by the actor by adding it to the tasks.
Note that we made a choice here to defer only the execution of the
body. Is this choice a good one? Suppose we wanted to also defer the
evaluation of the operands? We'll come back to these questions later.

Finally, we have to provide a scheduler. In this simple scheduler
there is a task queue for each individual actor and there is a
runnable queue for actors that have non-empty task queues. A task is
represented by a thunk.

```
(define runnable-actors)

(define (init-actors)
  (set! runnable-actors (queue:make)))
```

The procedure run-one either returns a task to be performed or it
signals that there are no tasks to be performed or it signals that
a runnable actor actually has no work to do.

```
(define (run-one)
  (if (queue:empty? runnable-actors)
      'nothing-to-do
      (let ((actor (queue:get-first! runnable-actors)))
        (if (MITscheme-continuation? actor)
            (lambda () (actor 'go))
            (let ((actor-task-queue (get-actor-task-queue actor)))
              (if (queue:empty? actor-task-queue)
                  (begin ;; Nothing for this actor to do.
                    (set-actor-runnable! actor #f)
                    'try-again)
                  (let ((task (queue:get-first! actor-task-queue)))
                    (if (queue:empty? actor-task-queue)
                        (set-actor-runnable! actor #f)
                        (queue:add-to-end! runnable-actors actor))
                    task)))))))
```

An actor that is already on the runnable-actors queue has its runnable
bit set. This allows the scheduler to avoid putting multiple
references to the same actor on the runnable actors queue.

If work is preempted by time sharing, the continuation of that work
is also put on the runnable-actors queue, and turned into a thunk, so
it can be executed as if it were a task.

The scheduler is the procedure below. It calls run-one to get work to
do. If it gets something to do, it executes it as a thunk and then
goes back for more. If it runs out, it returns. The procedure
atomically takes a thunk and executes it without allowing any
interrupts, using the without-interrupts feature of MIT/GNU Scheme.
Note that run-one is always executed atomically, preventing
reader-writer problems on the queues it manipulates.

```
(define (run)
  (let lp ((to-do (atomically run-one)))
    (cond ((eq? to-do 'try-again)
           (lp (atomically run-one)))
          ((eq? to-do 'nothing-to-do)
           to-do)
          (else (to-do)
                (root-continuation 'go)))))
```

Here we only enqueue an actor to be run if it is not already on the
runnable queue.

```
(define (add-to-runnable! actor)
  (atomically
   (lambda ()
     (if (actor-runnable? actor)
         'already-runnable
         (begin (set-actor-runnable! actor #t)
                (queue:add-to-end! runnable-actors actor)
                'made-runnable)))))
```

We can give an actor new tasks to perform. This always puts it on the
runnable queue.

```
(define (add-to-tasks! actor task)
  (atomically
      (lambda ()
        (queue:add-to-end! (get-actor-task-queue actor) task)))
  (add-to-runnable! actor)
  'task-added)
```


The read-eval-print loop has to be updated to correctly start the
system:

```
(define (init)
  (set! the-global-environment
      (extend-environment '() '() the-empty-environment))
  (init-actors)                                ;in schedule.scm
  (repl))

(define (repl)
  (if (eq? the-global-environment 'not-initialized)
      (error "Interpreter not initialized. Run (init) first."))
  (setup-time-sharing run)
  (let lp ((input (read)))
    (let ((output (eval input the-global-environment)))
      (write-line output))
    (lp (read))))

(define (go) (repl))
```

There is a preemptive timer interrupt that runs this whole mess.
There is an MIT/GNU Scheme specific detail. register-timer-event is
the MIT/GNU Scheme mechanism for delivering a timer interrupt. When
the time specified by its first argument expires, it invokes the
second argument.
This part of the system was very hard to get right: GJS was up all
night debugging it!

```scheme
(define (setup-time-sharing thunk)
  (set! user-continuation thunk)
  (call-with-current-continuation
   (lambda (k)
     (set! root-continuation k)
     (start-time-sharing)))
  ;; note: when time-sharing terminates the
  ;;user continuation is run again.
  (user-continuation))

(define time-sharing:quantum 100)
(define time-sharing-enabled? #t)
(define time-sharing? #f)
```

In fact, to help understand this, Pavel rewrote this procedure,
refactoring it to clarify the organization

```scheme
(define (start-time-sharing)
  (set! time-sharing? time-sharing-enabled?)
  (define (setup-interrupt)
    (if time-sharing?
        (register-timer-event time-sharing:quantum
                              on-interrupt)))
  (define (on-interrupt)
    (setup-interrupt)                ; Recursive interrupt setup
    (call-with-current-continuation
     (lambda (worker-continuation)
       (define (grab-a-continuation)
         (atomically
          (lambda ()
            (queue:add-to-end!
             runnable-actors worker-continuation))))
       (grab-a-continuation)
       (root-continuation 'go))))
  (setup-interrupt))

(define (stop-time-sharing)
  (set! time-sharing? #f))
```

There are a few more things that are necessary to make this system.
All user variables that are modified must be protected from
reader-writer problems.

```
(define (set-variable-value! var val env)
  (let plp ((env env))
    (if (eq? env the-empty-environment)
        (error "Unbound variable -- SET!" var)
        (let scan
            ((vars (vector-ref env 0))
             (vals (vector-ref env 1)))
          (cond ((null? vars) (plp (vector-ref env 2)))
                ((eq? var (car vars))
                 (double-check-lock
                   (lambda () (eq? var (car vars)))
                   (lambda ()
                     (set-car! vals val))
                   (lambda () (plp env))))
                (else (scan (cdr vars) (cdr vals))))))))
```

The assignment of the variable is guaranteed to be atomic. There is a
similar, but somewhat hairier problem with define-variable!.

Problem 9.1:

Here is a simple program that is implemented with actors.

```
(define (foo n)
  (define (buzz m)
    (if (not (= m 0)) (buzz (- m 1))))
  (define iter
    (alpha (l i)
       (if (not (= i 0))
           (begin
             (if (eq? l 'a) (buzz (* 100 i)) (buzz (* 100 (- n i))))
             (pp (list l i))
             (iter l (- i 1))))))
  (iter 'a n)
  (iter 'b n))
```

When I ran this program I got the following output:

```
eval> (foo 10)
actor-applied

eval> (b 10)
(b 9)
(b 8)
(a 10)
(b 7)
(b 6)
(a 9)
(b 5)
(b 4)
(a 8)
(b 3)
(a 7)
(a 6)
(b 2)
(a 5)
(b 1)
(a 4)
(a 3)
(a 2)
(a 1)
```

You may get a different order than I got. Why do we get this output in such a strange order? Explain what you see here.

Problem 9.2:

Note that we have made a choice in the handler for application of actor procedures: we defer only the execution of the body. Is this choice a good one? Suppose we wanted to also defer the evaluation of the operands?

a. Explain, in a short clear paragraph your opinion on this matter.

b. Write an apply handler for actors that defers the evaluation of operands. Demonstrate it. Does it have the properties you expect?

c. Would it be advantageous to supply another kind of object, perhaps defined by a "beta expression" that has this behavior?

Problem 9.3:

The procedure double-check-lock is used in set-variable-value! and
define-variable!.

```
(define (double-check-lock check do if-not)
  (let ((outside
          (atomically
           (lambda ()
             (if (check)
                 (begin (do)
                        (lambda () 'ok))
                 if-not)))))
    (outside)))
```

Why is it needed? Explain how it works.

Problem 9.4:

The first "actor implementation" of the Fibonacci procedure was

```
(define fib1
  (alpha (n c)
    (if (< n 2)
        (c n)
        (let ((x 'not-ready) (y 'not-ready))
          (define wait-for-xy
            (alpha (k)
                   (if (boolean/or (eq? x 'not-ready)
                                   (eq? y 'not-ready))
                       (wait-for-xy k)
                       (k #t))))
          (fib1 (- n 1) (lambda (v) (set! x v)))
          (fib1 (- n 2) (lambda (v) (set! y v)))
          (wait-for-xy (lambda (ignore) (c (+ x y)))))))))
```

a. Why is this so slow compared with the traditional version of fib
defined as a doubly-recursive function in the same interpreter?

Perhaps it is worthwhile to consider the following alternative
implementation of Fibonnaci:

```
(define fib2
  (alpha (n c)
    (if (< n 2)
        (c n)
        (let ((x 'not-ready) (y 'not-ready))
          (define check-if-done
            (lambda ()
              (if (boolean/or (eq? x 'not-ready)
                              (eq? y 'not-ready))
                  #f
                  (c (+ x y)))))
          (fib2 (- n 1)
                (lambda (v)
                  (set! x v)
                  (check-if-done)))
          (fib2 (- n 2)
                (lambda (v)
                  (set! y v)
                  (check-if-done)))))))
```

The essential difference is that the end test in the first version is
implemented as an actor loop, whereas in the second version it is
implemented as a Scheme procedure.

b. The fib2 implementation happens to have a subtle bug. Sometimes it
prints the result multiple times: try (fib2 10 write-line) a few
times, if you don't see ths problem (it is a heisenbug -- see
Wikipedia) try other numbers, like 8,9,11. Can you explain why this
occurs? How could you fix it? Caution: this is very hard.

Futures

The procedures that were illustrated in Problem 9.4 share an idea. An actor is started; when it finishes it sets a return value in the caller. The caller must check when the results of pending computations are completed for it to proceed to use those results. This pattern should be abstracted. One way is by the introduction of "future"s. (This abstraction was introduced by Robert H. Halstead, Jr. in the early 1980's. Sometimes futures are called "promises", but this word usually refers to delay thunks.) With futures, the code for Fibonacci looks like:

```
(define fib
  (alpha (n c)
    (if (< n 2)
        (c n)
        (let ((xp (future (lambda (k) (fib (- n 1) k))))
              (yp (future (lambda (k) (fib (- n 2) k)))))
          (wait xp
            (lambda (x)
              (wait yp
                (lambda (y)
                  (c (+ x y)))))))))))

eval> (fib 10 write-line)
actor-applied
55
```

Notice that the future converts the continuation into an object whose fulfillment one can wait for.

A future is a mutable data structure. It contains a continuation, a "done" flag, a value. When a future gets fulfilled, its value is stored, and the done flag is set.

Problem 9.5:

Here we implement futures. The implementation should be added to the guest interpreter.

a. Your first job is to construct the data structure for a future. We recommend using record structures for this.

b. Implement procedures FUTURE and WAIT. Make sure that your WAIT does not hang: it should return immediately but put a continuation on a queue of things to do. Assume here that a future is waited on only once. Make sure that you protect critical regions with ATOMICALLY.

c. What goes wrong with your implementation if the assumption is
    violated?

d. Improve your implementation to relax the assumption.